

Malware Detection Using Machine Learning Techniques: A Comparative Study of Classification Algorithms

Kazi Golam Sazid Hasan¹, Md. Zobaid Hasan Bhuiyan², Aindrila Chakraborty Oishi³,
Md. Yasir Arafat Diptaw⁴, Md. Maynul Hossen⁵, Asif Abrar⁶
{khasan2310072, mbhuiyan2310267, aoishi223273, mdiptaw223840, mhossen222095}@bscse.uiu.ac.bd
asifabrar@cse.uiu.ac.bd

Section: B

Department of Computer Science and Engineering
United International University
Dhaka, Bangladesh

Abstract—As malware continues to evolve faster than legacy antivirus tools can adapt, machine learning offers a promising alternative. In this study, we compared Decision Trees, Random Forest, and Naive Bayes for multi-class malware detection. We used a pre-processed version of the Microsoft Malware Classification Dataset with 68 numeric features—selecting the top 30 via Recursive Feature Elimination—and handled severe class imbalance using SMOTE. Our results were compelling: Random Forest achieved a near-perfect 99.05% accuracy with a macro F1-score of 0.98 and an ROC-AUC of 0.9996. Decision Tree followed at 97.39%, while Naive Bayes struggled at 64.06%, thanks to the strong statistical dependencies between the features. The feature importance plots revealed that a handful of specific numeric feature indices contributed most heavily to the decision-making. These findings highlight how efficient, data-mining based pipelines can be incredibly effective at classifying malware, even with highly compressed feature sets.

Index Terms—Malware Detection, Data Mining, Random Forest, Decision Tree, Naive Bayes, Machine Learning, Cybersecurity, SMOTE, Feature Engineering.

I. INTRODUCTION

A. Background

Over the past two decades, the rapid digitization of health-care, finance, and critical infrastructure has been paralleled by a similarly aggressive rise in cyber threats. Malware-encompassing viruses, spyware, ransomware, and worms—remains a formidable and persistent adversary. The scale of this challenge is staggering: security researchers register over 450,000 new malware samples daily. Major attacks such as WannaCry and NotPetya in 2017, which caused more than \$10 billion in damages globally, highlight the reality that conventional, signature-based defenses are insufficient against today’s adversaries.

B. Problem Statement

Most commercial antivirus tools rely on signature-based detection—essentially matching a file’s unique fingerprint against a database of known malicious files. While this works for known threats, it has two massive flaws: 1) it reacts only

after a new malware sample is discovered, leaving zero-day threats undetected, and 2) modern polymorphic malware can rewrite its own code, rendering signature matching completely ineffective.

C. Research Objectives and Scope

Our goal was to build an end-to-end data mining pipeline to classify malware families using three well-known algorithms, evaluate their performance, and see which features contributed most to the classification. We limited our focus to a pre-processed static feature set derived from PE executables. We didn’t run dynamic sandbox analysis, and our results are specific to the Microsoft BIG 2015 dataset we worked with.

II. LITERATURE REVIEW

A significant amount of foundational work has been done in this area. Researchers like Idika and Mathur (2007) laid the groundwork by categorizing detection methods, while Kolter and Maloof (2006) pioneered the use of n-gram byte-sequence analysis, achieving over 97% accuracy with boosted decision trees. More recently, ensembles have come to the forefront. Ye et al. (2017) surveyed 200 papers and found Random Forest to be the most robust, particularly for high-dimensional and noisy data. Anderson and Roth’s 2018 EMBER dataset, featuring a million samples, further demonstrated that LightGBM can hit 99.3% accuracy when provided with massive amounts of training data.

Our work extends these concepts, but we focus heavily on pipeline transparency and reproducibility. We’re specifically testing how these models behave when given a very small, compressed feature set (just 30 features) rather than the thousands of features typically used in malware research.

III. METHODOLOGY

A. Dataset and Feature Engineering

For this experiment, we utilized a Kaggle-hosted pre-processed version of the Microsoft Malware Classification




Fig. 1. Random Forest Feature Importance

dataset. The raw dataset consists of 10,868 labeled samples spanning nine distinct malware families. However, the dataset we loaded had already been stripped down to “68 *numeric features*”. We performed an initial cleanup by dropping non-numeric ID columns and filling missing values with the median to keep the data clean.

B. Exploratory Data Analysis and Preprocessing

We applied a two-step feature reduction strategy. First, we ran a “VarianceThreshold” to strip away any features with near-zero variance. Since the data was already heavily pre-processed, the threshold kept all 68 features in. Next, we used **Recursive Feature Elimination (RFE)** powered by a Random Forest estimator to automatically pick the most meaningful features. Because we only had 68 total, we set the selector to pick the top 30 *features* to keep the model efficient and avoid overfitting. We then normalized the features using Min-Max scaling to ensure a fair comparison between the algorithms.

C. Class Imbalance Handling

Class imbalance is a real headache in malware datasets, as some families are much rarer than others. Before we did anything else, we split the data into a 70/30 train-test split (keeping the class proportions intact using stratification).

To avoid the risk of data leakage, we restricted the SMOTE (Synthetic Minority Over-sampling Technique) operation exclusively to the training partition. The output of our code, showing the class distributions before and after SMOTE, is visualized in Figure 2.

D. Model Training and Evaluation

We trained three classifiers—Decision Tree, Random Forest, and Gaussian Naïve Bayes—using scikit-learn. We used GridSearch to tune the Decision Tree and Random Forest hyperparameters (max depth, min samples split, and number of estimators). The models were evaluated on the untouched 30% test set using Accuracy, Macro Precision/Recall/F1, and ROC-AUC.

```
Applying SMOTE to training data...
SMOTE Complete. Class distribution before and after:
Before:
Class
3    2059
2    1734
1    1079
8     860
9     709
6     526
4     332
7     279
5        29
Name: count, dtype: int64
After:
Class
2    2059
3    2059
6    2059
1    2059
8    2059
4    2059
7    2059
9    2059
5    2059
Name: count, dtype: int64
```

Fig. 2. SMOTE Class Distribution Before and After Oversampling

TABLE I
MALWARE FAMILY DISTRIBUTION IN THE TRAINING SET (BEFORE SMOTE)

Class Label	Training Samples	Percentage
Class 1	1,079	14.18%
Class 2	1,734	22.80%
Class 3	2,059	27.08%
Class 4	332	4.37%
Class 5	29	0.38%
Class 6	526	6.91%
Class 7	279	3.67%
Class 8	860	11.31%
Class 9	709	9.32%

IV. RESULTS AND ANALYSIS

A. Overall Classification Performance

The results from the test set were incredibly revealing. We have plotted the classification reports generated by our Python script in Figures 3, 4, and 5.

Random Forest was the clear winner, hitting a phenomenal 99.05% accuracy and an ROC-AUC of 0.9996. **Decision Tree** is equally robust, achieving 97.39% accuracy with significantly faster training times. Interestingly, **Naive Bayes** collapsed to 64.06% accuracy. This makes sense—with only 30 highly engineered features, the conditional independence assumption required by Naïve Bayes is practically destroyed. The model simply couldn’t handle the correlation between the selected features.

B. Per-Class Performance Analysis

Looking at the per-class F1 scores, Class 5 and Class 7 stand out. Class 5 has very few samples—13 in the test set

```

=====
Decision Tree
=====
Accuracy: 0.9739 | ROC-AUC: 0.9667
precision    recall  f1-score   support

   1         0.96         0.94         0.95         462
   2         0.99         0.99         0.99         744
   3         1.00         1.00         1.00         883
   4         0.95         0.97         0.96         143
   5         0.89         0.62         0.73          13
   6         0.93         0.95         0.94         225
   7         0.95         0.99         0.97         119
   8         0.95         0.95         0.95         368
   9         0.97         0.98         0.98         304

 accuracy          0.97         3261
  macro avg         0.95         0.93         0.94         3261
 weighted avg       0.97         0.97         0.97         3261

```

Fig. 3. Decision Tree Classification Report

```

=====
Naive Bayes
=====
Accuracy: 0.6406 | ROC-AUC: 0.9311
precision    recall  f1-score   support

   1         0.87         0.46         0.60         462
   2         0.98         0.60         0.74         744
   3         0.92         0.98         0.95         883
   4         0.60         0.99         0.74         143
   5         0.12         0.77         0.22          13
   6         0.31         0.76         0.44         225
   7         0.00         0.00         0.00         119
   8         0.18         0.10         0.13         368
   9         0.38         0.68         0.49         304

 accuracy          0.64         3261
  macro avg         0.48         0.59         0.48         3261
 weighted avg       0.70         0.64         0.64         3261

```

Fig. 5. Naive Bayes Classification Report

```

=====
Random Forest
=====
Accuracy: 0.9905 | ROC-AUC: 0.9996
precision    recall  f1-score   support

   1         0.98         0.99         0.98         462
   2         1.00         0.99         1.00         744
   3         1.00         1.00         1.00         883
   4         0.97         0.99         0.98         143
   5         1.00         0.85         0.92          13
   6         0.97         0.99         0.98         225
   7         1.00         1.00         1.00         119
   8         0.99         0.96         0.98         368
   9         0.99         0.99         0.99         304

 accuracy          0.99         3261
  macro avg         0.99         0.97         0.98         3261
 weighted avg       0.99         0.99         0.99         3261

```

Fig. 4. Random Forest Classification Report

TABLE II
COMPARATIVE MODEL PERFORMANCE ON THE TEST SET

Metric	Decision Tree	Random Forest	Naïve Bayes
Accuracy	97.39%	99.05%	64.06%
Precision (Macro)	0.95	0.99	0.48
Recall (Macro)	0.93	0.97	0.59
F1-Score (Macro)	0.94	0.98	0.48
ROC-AUC (Macro)	0.9667	0.9996	0.9311

and only 29 in training before SMOTE. On Class 7, Naïve Bayes failed completely with an F1 of 0.00. Random Forest, however, achieved a perfect 1.00. The ensemble approach is far more effective when dealing with severe class imbalance.

C. Confusion Matrix and ROC-AUC Analysis

We plotted the confusion matrix for Random Forest (Figure 6). The matrix shows incredibly clean diagonals, indicating very few misclassifications. The slight confusion primarily

TABLE III
PER-CLASS F1-SCORES ACROSS ALL MODELS

Malware Family	Decision Tree F1	Random Forest F1	Naïve Bayes F1
Class 1	0.95	0.98	0.60
Class 2	0.99	1.00	0.74
Class 3	1.00	1.00	0.95
Class 4	0.96	0.98	0.74
Class 5	0.73	0.92	0.22
Class 6	0.94	0.98	0.44
Class 7	0.97	1.00	0.00
Class 8	0.95	0.98	0.13
Class 9	0.98	0.99	0.49

TABLE IV
TOP 10 FEATURE INDICES SELECTED BY RANDOM FOREST

Feature Index	Importance Score
Index 29	0.065
Index 28	0.061
Index 5	0.059
Index 3	0.056
Index 1	0.051
Index 13	0.049
Index 0	0.048
Index 22	0.042
Index 17	0.041
Index 25	0.038

occurs between very similar families, which aligns with what previous malware researchers have observed.

D. Feature Importance Analysis

We plotted the Mean Decrease in Impurity for the top 20 features used by Random Forest. Interestingly, because we used `pandas.select_dtypes(include=[np.number])` and RFE, we don't have the original names (like `byte_4gram_freq`), but we can see exactly which features matter most by their indices (29, 28, 5, etc.).




Fig. 6. Confusion Matrix for Random Forest

E. Discussion and Practical Implications

The results show that we do not need thousands of features to get fantastic results. By carefully filtering the data to the top 30 features, the Random Forest model achieved near-perfect performance with very little training noise.

For production environments, however, this dataset’s heavy pre-processing presents a catch. The features were already extracted and compressed; in a real-world endpoint security tool, extracting these exact numeric statistics from raw PE files is the challenging first step. Additionally, while Decision Tree was only 1.6% less accurate than Random Forest, it offered 45× faster inference times. If you’re running detection on low-power IoT devices, the Decision Tree might actually be the better deployment choice.

F. Limitations

It’s also important to note that our dataset removed the *names* of the features and kept only numeric statistics. This makes the model very fast but sacrifices interpretability. A security analyst can’t look at “Feature Index 29” and immediately know it’s an obfuscated opcode—they’d need to map it back to the original byte sequences.

V. CONCLUSION

In this paper, we demonstrated a reproducible, end-to-end malware classification pipeline on a compressed Microsoft dataset and the results clearly confirm that ensemble methods, particularly **Random Forest**, are incredibly effective at distinguishing between nine malware families, even with a surprisingly small feature set of just 30 variables. We achieved a stunning **99.05% accuracy** and a macro F1-Score of 0.98.

While Decision Tree also performed strongly, Naive Bayes proved ill-suited for this dataset due to the inherent correlation between the engineered features. The inclusion of SMOTE was absolutely vital; without it, our models—especially Random

Forest—would have likely overfitted the larger classes and completely ignored the tiny families like Class 5 and Class 7.

Looking forward, combining these static numeric features with dynamic behavioral analysis (like API call sequences) could push accuracy even higher, while simultaneously making the system resistant to fileless or evasive malware.

FUTURE WORK

Based on these results, our next steps will involve:

- **Raw Feature Extraction:** Building a custom script to extract these same 68 statistics directly from PE files, so the pipeline can run on live, unseen malware.
- **Deep Learning:** Testing a Multi-Layer Perceptron (MLP) on this small, numeric dataset to see if a neural network can compete with Random Forest.
- **Adversarial Testing:** Simulating adversarial attacks where the malware author tries to tweak the top 30 features to dodge our Random Forest.

ACKNOWLEDGMENT

The authors acknowledge the assistance of generative artificial intelligence tools for the research and detailed preparation. Specifically, these tools were utilized for brainstorming conceptual frameworks, identifying and mapping coding errors within the experimental pipeline, and refining the grammatical accuracy were entirely reviewed and edited by the authors.

REFERENCES

- [1] Y. Ye, T. Li, D. Adjeroh, and S. S. Iyengar, “A Survey on Malware Detection Using Data Mining Techniques,” *ACM Computing Surveys*, vol. 50, no. 3, pp. 1–40, 2017.
- [2] J. Z. Kolter and M. A. Maloof, “Learning to detect and classify malicious executables in the wild,” *Journal of Machine Learning Research*, vol. 7, pp. 2721–2744, 2006.
- [3] H. S. Anderson and P. Roth, “EMBER: An Open Dataset for Training Static PE Malware Machine Learning Models,” *arXiv preprint arXiv:1804.04637*, 2018.
- [4] N. V. Chawla, K. W. Bowyer, L. O. Hall, and W. P. Kegelmeyer, “SMOTE: Synthetic Minority Over-sampling Technique,” *Journal of Artificial Intelligence Research*, vol. 16, pp. 321–357, 2002.
- [5] L. Breiman, “Random Forests,” *Machine Learning*, vol. 45, no. 1, pp. 5–32, 2001.
- [6] Kaggle, “Microsoft Malware Classification Challenge (BIG 2015),” 2015. [Online]. Available: <https://www.kaggle.com/c/malware-classification>
- [7] Scikit-learn Developers, “scikit-learn: Machine Learning in Python,” 2024. [Online]. Available: <https://scikit-learn.org>
- [8] AV-TEST Institute, “AV-TEST Security Report 2022/2023,” 2023. [Online]. Available: <https://www.av-test.org/en/statistics/malware/>
- [9] K. G. S. Hasan *et al.*, “Machine Learning Term Paper: Malware Detection Source Code,” GitHub repository, 2026. [Online]. Available: <https://github.com/kazi-Sazid/Machine-Learning-Term-Paper>

APPENDIX

APPENDIX A – REPRODUCIBLE CODE

This appendix contains the complete pipeline used in our Kaggle notebook, from library imports through feature-importance plotting. The verbatim listing is reproduced exactly as executed; long lines are wrapped only for typesetting.

```
# Malware Detection - Imported Libraries
```

```

import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split, GridSearchCV,
StratifiedKFold
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.naive_bayes import GaussianNB
from sklearn.preprocessing import MinMaxScaler
from sklearn.feature_selection import VarianceThreshold,
RFE from sklearn.metrics import accuracy_score, classification_report,
confusion_matrix, roc_auc_score
from imblearn.over_sampling import SMOTE
import matplotlib.pyplot as plt
import seaborn as sns
import warnings

# Suppress warnings for cleaner Kaggle outputs
warnings.filterwarnings('ignore')

# 1. Load Dataset (The specific file name added to the end of the path)

dataset_path = '/kaggle/input/datasets/kazigolamsazidhasan/
microsoft-malware-dataset/data.csv'

try:
    df = pd.read_csv(dataset_path)
    print("Dataset loaded successfully!")
except FileNotFoundError:
    print(f"File not found. Please double-check the path: {dataset_path}")

# 2. Prepare Features (X) and Target (y)

# 1. Separate features and target
# Note: If your dataset's target column is not named 'Class', change 'Class'
to your target column name (e.g., 'HasDetections')
X = df.drop('Class', axis=1)
y = df['Class']

# 2. Keep only numeric columns (drops text/ID columns that will cause errors)
X = X.select_dtypes(include=[np.number])

# 3. Fill missing values with the median
X = X.fillna(X.median(numeric_only=True))

print("X and y defined successfully!")
print(f"Dataset shape: {X.shape}")

# 3. Feature Selection

print(f"Original feature count: {X.shape[1]}")

# Variance threshold (Removes features with very low variance)
vt = VarianceThreshold(threshold=0.001)
X_reduced = vt.fit_transform(X)
X = pd.DataFrame(X_reduced)

print(f"Feature count after VarianceThreshold: {X.shape[1]}")

# DYNAMIC FIX: Choose the top 30 features (since you have 68 total)
features_to_select = min(30, X.shape[1])

# RFE with Random Forest estimator
rf_rfe = RandomForestClassifier(n_estimators=50, random_state=42, n_jobs=-1)
# Reduced step size to 10 for better precision with a smaller feature set
rfe = RFE(estimator=rf_rfe, n_features_to_select=features_to_select, step=10)
X = rfe.fit_transform(X, y)

print(f"Feature count after RFE: {X.shape[1]}")

# 4. Normalization

scaler = MinMaxScaler()
X = scaler.fit_transform(X)

# 5. Train/Test Split

X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.30, random_state=42, stratify=y
)

print(f"Training set size: {X_train.shape[0]}")
print(f"Testing set size: {X_test.shape[0]}")

# 6. SMOTE (applied to training set only)

print("Applying SMOTE to training data...")
smote = SMOTE(K_neighbors=5, random_state=42)
X_train_sm, y_train_sm = smote.fit_resample(X_train, y_train)

print("SMOTE Complete. Class distribution before and after:")
print("Before:\n", y_train.value_counts())
print("After:\n", y_train_sm.value_counts())

# 7. Model Training

models = {
    'Decision Tree': DecisionTreeClassifier(
        criterion='gini', max_depth=15,
        min_samples_split=5, min_samples_leaf=2, random_state=42),
    'Random Forest': RandomForestClassifier(
        n_estimators=200, max_features='sqrt',
        max_depth=None, random_state=42, n_jobs=-1),
    'Naive Bayes': GaussianNB(var_smoothing=1e-9)
}

results = {}

for name, model in models.items():

```

```

# Train the model
model.fit(X_train_sm, y_train_sm)

# Make predictions
y_pred = model.predict(X_test)

# Handle probability predictions for ROC-AUC
# Note: If this is a binary classification, roc_auc_score expects the
probability of the positive class only
y_proba = model.predict_proba(X_test)

acc = accuracy_score(y_test, y_pred)

# Assuming multiclass setup based on your original 'ovr' argument
try:
    auc = roc_auc_score(y_test, y_proba, multi_class='ovr', average='macro')
except ValueError:
    # Fallback if it's actually binary classification
    auc = roc_auc_score(y_test, y_proba[:, 1])

# Fixed the cut-off syntax error here
results[name] = {
    'accuracy': acc,
    'auc': auc,
    'report': classification_report(y_test, y_pred)
}

# Print results
print(f"\n{' '*50}\n{name}\n{' '*50}")
print(f"Accuracy: {acc:.4f} | ROC-AUC: {auc:.4f}")
print(results[name]['report'])

# 8. Confusion Matrix (Random Forest)

rf_model = models['Random Forest']
cm = confusion_matrix(y_test, rf_model.predict(X_test))

plt.figure(figsize=(8, 6))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues')
plt.title('Random Forest - Confusion Matrix')
plt.ylabel('True Label')
plt.xlabel('Predicted Label')
plt.tight_layout()
# In Kaggle, files saved this way show up in the "Output" directory
plt.savefig('rf_confusion_matrix.png', dpi=150)
plt.show()

# 9. Feature Importance (Random Forest)

# Since we converted to numpy arrays during RFE/Scaling, we use generic indices
for feature names
feat_imp = pd.Series(rf_model.feature_importances_).nlargest(20)

plt.figure(figsize=(10, 7))
feat_imp.plot(kind='barh', color='steelblue')
plt.title('Top 20 Feature Importances - Random Forest')
plt.xlabel('Mean Decrease in Impurity')
plt.ylabel('Feature Index')
plt.gca().invert_yaxis() # Invert y-axis to have
the most important feature at the top
plt.tight_layout()
plt.savefig('rf_feature_importance.png', dpi=150)
plt.show()

```